

Breaking the Cube

A gentle introduction

Warmin'

Trollin'

Profit

Warmin

Don't worry I'll run it in a VM!

Virtualization

“Virtualization, in computing, refers to the act of creating a virtual (rather than actual) version of something, including but not limited to a virtual computer hardware platform, operating system (OS), storage device, or computer network resources.”

2 types of virtualization

- Hardware-based Virtual Machine (HVM)
- Para-virtualization

HVM

- Uses Intel VT-x in order to provide isolation
- Introduces a “new” ring mode, un/privileged mode
- Privileged instructions must be emulated by the Hypervisor
 - Great attack surface, parsing x86 is tedious
- Memory mode has to be decided by VMM
 - Huge headache to code
- Device security must be enumerated
 - Attacks against VT-d, etc

Paravirtualization

- First introduced by Xen
- Modify the Guest OS and remove all privileged instructions
- Kinda better performance than HVM due to MMU
- In Linux mainline since 2.6.23
- Hyper-V calls it “enlightened VM”
- During boot of a Windows guest machine it can detect it is virtualized by Hyper-V and thus become “enlightened”

Guest-Host comm

- How does a VM asks for more memory ?
- How does a VM perform I/O operations?
- How does a VM perform a context switch?

Guest-Host comm

- HVMs offer a “trap-based” interface
 - Implemented as a bitarray of events in the Hypervisors
 - Later Intel + AMD offered “Virtualization exceptions” for custom operations
- Paravirtualization offers hypercalls
 - Special address range to jump into to cause a trap


Rollin'

2 main types of escapes

- Golden Escapes
 - Direct escape from an unprivileged VM to the hypervisor
- Chained escapes
 - Exploitation of several bugs in order to reach the hypervisor
 - May it be logic or a memory corruption

Golden Escapes

- Direct escape from Guest to the Hypervisor
- Either by faulty hypercalls or by bad interpretation of an event
- Use your imagination
- These are GOLDEN escapes
- I only know one public



XSA-7

Golden Escapes

- Kinda famous one
- Originally found somewhere in 2k6 (CVE-2006-044)
- Caused by improper understanding of the #GP
 - Intel said one thing, AMD said another
 - Xen only implemented AMD's for both CPUs
- Resulted in a golden escape from DomU <-> Hypervisor

Golden Escapes

- On 64b only 48bits are used for the address space
 - This gives us only a maximum of 256 terabyte of RAM :(
- If you look at a typical 64b address you'd see that bits 48-63 are all the same as bit 47
 - For ex `0xffff8000deadbeef`
- Intel made this on purpose to stop from programmers from using those bits as special flags

Golden Escapes

- Intel called this range of valid addresses “canonical”
 - Meaning we can use it
 - 0x0000000000000000-0x00007fffffffffff
 - 0xffff800000000000-0xffffffffffffffff
- What happens if we try use (execute from) a non-canonical address then?

Listen Tight

A #GP fault occurs

Golden Escapes

- But this is where the magic comes
- On AMD the #GP occurs while the machine is in **guest mode**
- On Intel the #GP occurs while the machine is in **hypervisor mode**

Golden Escapes


- While the #GP occurs, Xen restores all of it's registers from the stack
- If we're in Hypervisor mode, we'd get the **Hypervisor stack**
- But...
- If we're in Guest mode, we'd get the **Guest stack**

That's Gold

Golden Escapes

- Allocate a page in a non-canonical address
- Create simple shellcode in a non-canonical-address-1 (valid one)
 - So when we return from the syscall we'd generate a #GP
- Setup fake stack that the hypervisor will restore from
- Profit

Chained Escapes



Chained Escapes

- Kinda main topic :(
- I'd love to find more golden escapes :X
- They're usually composed by 2-3 bugs
 - Infoleak – read sensitive struct info from the hypervisor
 - Corruption or a logic bug on guest <-> dom0
 - Corruption or a logic bug on dom0 <-> hypervisor
 - Excluding driver domains as qemu is cheating

XSA-105

Chained Escapes

- HVM needs to be able to emulate privileged instructions
- LIDT, LGDT, INVLPG, LMSW
- The code which handles it resides in a huge case
- What could go wrong ?

Chained Escapes

- LIDT = Load a new Interrupt Descriptor Table
- LGDT = Load a new Global Descriptor Table
- INVPLG = Cache flush, invalidate the page cache
- LMSW = Overwrite CR0

All of these are
sensitive instructions

Parsing x86 is easy
right ?

Chained Escapes

```
3711 case 2: /* lgdt */
3712 case 3: /* lidt */
3713     generate_exception_if(ea.type != OP_MEM, EXC_UD, -1);
3714     fail_if(ops->write_segment == NULL);
3715     memset(&reg, 0, sizeof(reg));
3716     if ( (rc = read_ulong(ea.mem.seg, ea.mem.off+0,
3717                          &limit, 2, ctxt, ops)) ||
3718          (rc = read_ulong(ea.mem.seg, ea.mem.off+2,
3719                          &base, mode_64bit() ? 8 : 4, ctxt, ops)) )
3720         goto done;
3721     reg.base = base;
3722     reg.limit = limit;
3723     if ( op_bytes == 2 )
3724         reg.base &= 0xffffffff;
3725     if ( (rc = ops->write_segment((modrm_reg & 1) ?
3726                                x86_seg_idtr : x86_seg_gdtr,
3727                                &reg, ctxt)) )
3728         goto done;
3729     break;
```


No privilege checks!

Chained Escapes

1. Kinda generic attaq:
 1. Create a custom IDT table in usermode
 2. Cause an invalid opcode to be generated
 3. Replace the opcode to be LIDT
 1. Be quick enough to do so
2. Profit

XSA-84

```
79 static int flask_copyin_string(XEN_GUEST_HANDLE(char) u_buf, char **buf,  
80                               size_t size)  
81 {  
82     char *tmp = xmalloc_bytes(size + 1);  
83  
84     if ( !tmp )  
85         return -ENOMEM;  
86  
87     if ( copy_from_guest(tmp, u_buf, size) )  
88     {  
89         xfree(tmp);  
90         return -EFAULT;  
91     }  
92     tmp[size] = 0;  
93  
94     *buf = tmp;  
95     return 0;  
96 }  
97
```

Chained Escapes

- Kinda golden escape, but not a full one
- Finalizing this one needs another vulnerability
 - Other than an infoleak
 - And tmem internals..
- However this code is accessible directly using a hypercall
 - FLASK_SETBOOL hypercall

Chained Escapes

1. Try to allocate 4gb of memory
2. Xmalloc allocation add +1 by default
3. Profit from zero allocation

Profit

Profit

- I'd like to take things a bit to the next stage
- INSERT_HERE_PROGRAM_ANAL_BUZZ
- Below are some ideas I would like to try in the next several months

Anal

- I'd like to try to get Xen with Address Sanitizer
- This means modifying the TMEM implementation
- Sounds easy? It depends

From the xen-devel-ml:

“Until TMEM has gained production maturity, the Xen.org security team intends [..] to handle these and future TMEM vulnerabilities in public, as if they were normal non-security-related bugs.

We therefore intend that currently-known vulnerabilities will be publicly disclosed on the xen-devel mailing list, as normal bug reports, at the expiry of the XSA-15 embargo. In the meantime the list below may be helpful. “

Anal

- I'd like to get the Xen project compiled with Clang and have the ability to search for paths in it
- Maybe even use joern on it and model a few bugs
- I used ctags + cscope + grep in my research, until learning the whole tree to find the appropriate calls
 - Lots of indirect struct initialization, too many defines
- I'd like to be able to get basic graphing abilities
 - I'm using IDA originally, with my own compiled Xen
- Doesn't work for Guest <-> Host stuff..Lots of headaches

Anal

- Finally, I'd like to be able to set constraints on Xen and ask some solver to solve them for me
- Z3, yachs, stp, whatever
- Kinda gets all previous ideas into one project...
- Help is welcomed! :p I can share 0days

Thanks for listening

Questions?